

Introduction to SQL and SQLite3

Dr. John Artz

Outline

- What is a (Relational) Database
- Introduction to SQLite3 and Demo
- Creating Tables in SQL
- SQL Data Manipulation Statements
- Defining the SoupSales Tables
- Some Simple SQL Queries

What Is a Database?

- A Database Is An Organized Collection Of Data
 - Wikipedia
- Notes
 - Organization Implies Purpose
 - Collection Implies Containment
 - Data Implies Discreet Pieces of Information

What Is a Relational Database?

- A relational database is a digital database whose organization is based on the relational model of data, as proposed by E.F. Codd in 1970. This model organizes data into one or more tables (or "relations") of rows and columns, with a unique key for each row.
 - Wikipedia

A Relational Database is all Tables

Trade	Qty	Storeid	Shupld	Promoid	Sales
1	1	1	1	1	485.18
2	3	1	2	1	424.9
3	1	1	3	1	888.81
4	1	1	4	1	425.81
5	1	1	5	1	387.46
6	1	1	6	1	719.81
7	1	1	7	1	287.84
8	1	1	8	1	185.82

Storeid	Location	Elevatio	Size	Mgrid
1	Rockville	Basic	Medium	1
2	Potomac	Upscale	Medium	2
3	Gaithersburg	Basic	Large	3
4	Bethesda	Upscale	Medium	4
5	Fredrick	Basic	Large	5

Promoid	Medium	Target	Interval
1	Radio	Morning Drive	Weekday
2	Radio	Evening Drive	Weekday
3	Radio	Chime	Weekend
4	Newspaper	Coupon	Weekday
5	Newspaper	Insert	Weekday
6	Newspaper	Coupon	Weekend
7	Newspaper	Full Page Ad	Weekday
8	Newspaper	Coupon	Weekend

And Not Just Any Tables

Tables Must Be
Properly Designed



But, We Will Get
To That Later



Poorly
Designed
Tables Are
Worse Than
Useless

They Provide Incorrect Information

Tabular Data in General

- Tabular Data is Data Organized in Rows and Columns
- Most Data is Inherently Tabular Including: Spreadsheets, Cross Tabulations, Data Cubes, Sparse Tabular Data (e.g. NoSQL, XML), in Fact, Any Data Organized in Rows and Columns
- Relational Data is Tabular as Well

Tabular vs Relational

- All Relational Data is Tabular
- Not All Tabular Data is Relational
- In Fact, Most Tabular Data is Not Relational
- Relational Tables Must Be Designed Following a Rigorous Process Called Relational Database Design

Gloss on Relational DB Design

- Each Table Corresponds to a WELL DEFINED Entity Class or Category
- Each Row Corresponds to an Instance of That Category
- Each Fact in a Row is a Fact About That Instance
- Integrity Rules Prevent Corruption of the Data and Its Derivations

How Are Tables Designed Correctly?

- Conceptual Database Design Helps Determine if Categories are Well Formed
- Logical Database Design Helps Determine if Tables are Well Formed
- Integrity Rules Help Insure That Queries Produce the Expected Results

An Example

Categories (Relations)
Must be Well Defined

Row or Table
Design Has
Rules Too

Soups	Type	Vendor	Mode	Style
1	Chicken Noodle	Progresso	Canned	Basic
2	Chicken Noodle	Campbells	Canned	Basic
3	Chicken Noodle	Lipton	Dry	Basic
4	Chicken Noodle	Campbells	Canned	Chunky
5	Chicken Noodle	Wolfgang Puck	Canned	Gourmet
6	Chicken Noodle	Pacific Organic	Boxed	Organic
7	Chicken Noodle	Healthy Choice	Canned	Healthy
8	Split Pea	Campbells	Canned	Basic

Integrity Rules Require Rows to be the Same Kind of Thing And Address Relationships Between Categories

Relational vs. Not Relational (XML)

Relational

Soups
1
2
3
4
5
6
7
8

This is About a
Can of Soup

It is Unclear What
This is About


Not Relational (Now NoSQL)

```

<Soup>
<Id>1</Id>
<Type>Chicken Noodle</Type>
<Vendor>Progresso</Vendor>
  <Mode>Canned</Mode>
</Vendor>Campbells</Vendor>
  <Mode>Canned</Mode>
</Vendor>Lipton</Vendor>
  <Mode>Dried</Mode>
</Soup>

```

Relational vs. NoSQL (Cassandra)



Note: Each Row Can Have Different Attributes
Data is Accessible via Host Language Program
How Difficult Would a Query Be?

A Relational Table

Table Name → **Soups**

Unique Identifier → **ID**

Columns Are The Same for All Rows

Rows →

ID	Type	Vendor	Mode	Style
5	Chicken Noodle	Progresso	Canned	Basic
6	Chicken Noodle	Campbells	Canned	Basic
7	Chicken Noodle	Lipton	Dry	Basic
8	Chicken Noodle	Campbells	Canned	Chunky
9	Chicken Noodle	Wolfgang Puck	Canned	Gourmet
10	Minestrone	Pacific Organic	Boxed	Organic
11	Minestrone	Progresso	Canned	Basic
12	Minestrone	Campbells	Canned	Basic
13	Minestrone	Amy's Organic	Canned	Organic

Data is Accessible Using SQL
Host Language Programs Are Not Necessary

It Looks a Lot Like a Spreadsheet

- A Relational Table Does Look a Lot Like a Spreadsheet
- But, There is No Design Theory for Spreadsheets
- And Spreadsheet Technology Couldn't Handle it if There Was
- But, There is Design Theory for Relational Tables
- This Makes All the Difference in the World

Relational Database Snobbery


Only This is
a Database

SQL



These Are Not
They Are Just for Storage

Key-value



Document



Relational Terminology

- We Call a Table a Relation
- We Call a Column an Attribute
- We Call a Row a Tuple
- We Call the Unique Identifier, a Primary Key

A Relational Table, Revisited

Relation Name: **Soops**

Primary Key: **SoupId**

Attributes: **Type, Vendor, Mode, Style**


Tuple: (Row)

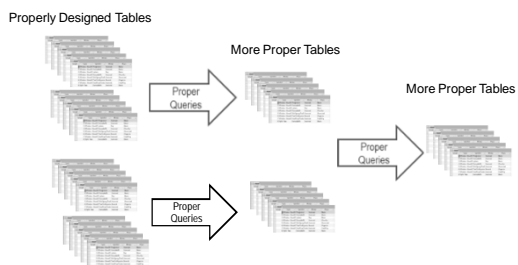
SoupId	Type	Vendor	Mode	Style
1	Chicken Noodle	Progresso	Canned	Basic
5	Chicken Noodle	Campbells	Canned	Basic
6	Chicken Noodle	Lipton	Dry	Basic
7	Chicken Noodle	Campbells	Canned	Chunky
8	Chicken Noodle	Woonfong Park	Canned	Gourmet
9	Chicken Noodle	Pacific Organic	Boxed	Organic
10	Minestrone	Progresso	Canned	Basic
11	Minestrone	Campbells	Canned	Basic
12	Minestrone	Amy's Organic	Canned	Organic

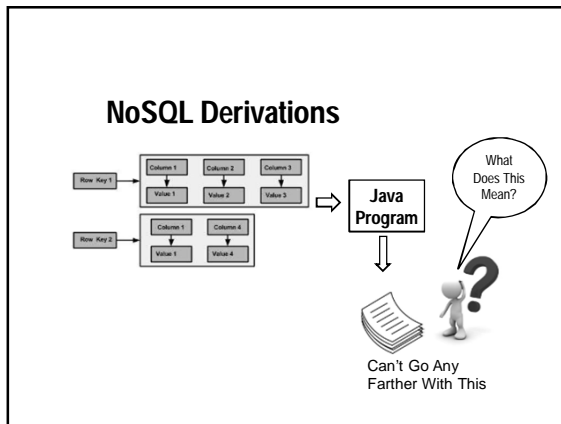
Attributes are Facts About the Thing
Identified By the Primary Key

Relational Design Restrictions

- All Rows Must Have the Same Attributes
- All Rows Must Have Unique Identifiers
- All Rows Must Be Instances of the Same Thing
- But, If You Follow All the Rules, You Can Derive New Information Using a High Level Query Language Such as Structured Query Language, or SQL


Relationally Derived Information

And It Doesn't Stop There





Relational Performance Problems


- Relational Databases Must Maintain Data Integrity in Order Answer Queries Correctly
 - Data Value Constraints
 - Concurrency Control
 - Transaction Models
- This Makes Relational Databases (Currently) Unacceptable for Big Data Applications




NoSQL Databases


cassandra


Couchbase


mongoDB


HBase

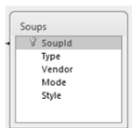
Big Data Has Given Rise To Numerous NoSQL Databases

Relational DBs Vs. NoSQL

- Relational Databases Are For the Derivation and Delivery of Information
 - They Focus on Design and Exploitation Discipline
- NoSQL Databases Are For the Storage and Retrieval of Data
 - They Focus on Storage and Retrieval Performance

Again: Relational vs. Not Relational

Relational



Structured Data

Not Relational

```

<Soup>
<Id>1</Id>
<Type>Chicken Noodle</Type>
<Vendor>Progresso</Vendor>
  <Mode>Canned</Mode>
<Vendor>Campbells</Vendor>
  <Mode>Canned</Mode>
<Vendor>Lipton</Vendor>
  <Mode>Dried</Mode>
</Soup>


```

Semi- Structured Data

Why Relational Snobbery?

Only This is a Database

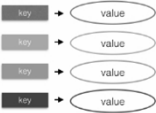
SQL




This is For Information

These Are Not

Key-value



Document




These Are For Storage

Why NoSQL Snobbery?

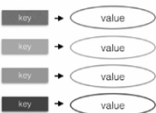
This is Slow, Restrictive And Low Volume

SQL




These Are Fast, Flexible and High Volume

Key-value



Document



Soon, This Will All Change



NewSQL

NewSQL

- The Goal of NewSQL is to Support Both High Volume Data and Traditional Relational Databases in a Single Platform
- In a Sound Byte, NewSQL = SQL + NoSQL
- This Will Be Achieved By Changing the Internal Architecture of Relational Databases
- SQL Will Probably Not Change That Much

What is SQL?

- SQL (Structured English Query Language) is the Standard Query Language for Relational Databases
- While Not All Relational Database Management Systems Stick to the Standard, Standard SQL Will Work With Any Relational DBMS

We Will Be Using SQLite

- "SQLite is a software library that implements a self-contained, serverless, zero-configuration, transactional SQL database engine. SQLite is the most widely deployed database engine in the world. The source code for SQLite is in the public domain."
 - Source: <https://www.sqlite.org/>

Why SQLite?

- SQLite is a software library that implements a self-contained, server-less, zero-configuration, transactional SQL database engine. SQLite is the most widely deployed database engine in the world. The source code for SQLite is in the public domain.

As a Practical Matter....

- SQLite3 is a Free Download
- It Does Not Require Installation
- It Handles Everything We Need To Do
 - SQLite3 is a Standalone Relational Database
 - PySQLite3 Allows Us to Interface with SQLite3 From a Python Program

To Avoid Confusion...

- SQL is a Data Language That Can Be Implemented in Any Relational Database Management System (RDBMS)
- SQLite3 is One of Many, Many RDBMS's

However, SQLite3 Doesn't Support

- High Transaction Rates
- Extremely Large Databases
- Access Control
- Client/Server
- Replication
- GUI Interface
- But, It Is Fast, Free, and Easy

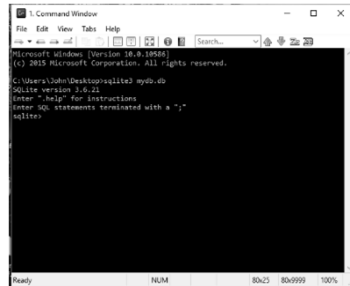
How Do You Get SQLite?

- Go to <https://www.sqlite.org/>
- Click on Download in the Top Menu Bar
- Select the Precompiled Binary for Your Machine
- Download it
- No Installation is Necessary

Some Tips

- I Created a Folder on My Desktop and Did Everything From Within That Folder
- This Just Keeps All My Stuff in One Place and Eliminates the Need for Path Names
- From Within That Folder Run a Cmd Shell
- In the Cmd Shell Type "SQLite3 MyDB.db"

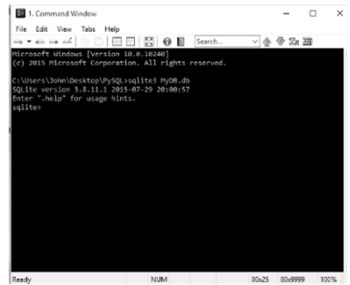
From the Cmd Shell



```
Microsoft Windows [Version 10.0.10586]
(c) 2015 Microsoft Corporation. All rights reserved.

C:\Users\John\Desktop\sqlite3>mydb.db
SQLite version 3.9.2
Enter \help for instructions
Enter SQL statements terminated with a \>
sqlite>
```

The SQLite Prompt



```
Microsoft Windows [Version 10.0.10586]
(c) 2015 Microsoft Corporation. All rights reserved.

C:\Users\John\Desktop\PySQL\sqlite3>mydb.db
SQLite version 3.9.2
Enter \help for instructions
Enter SQL statements terminated with a \>
sqlite>
```

A Quick SQLite3 Demo

- Invoke SQLite3
- Create a Table
- Verify the Table
- Insert a Row
- Verify the Row
- Update the Row
- Verify the Update

Some Handy dot Commands

- .help – list dot commands
- .tables – list tables
- .exit or .quit – close the SQLite3 Window
- .read *filename* – read SQL or SQLite3 commands from *filename*
- .import *filename table-name* – Import data from *filename* into *table-name*

SQL Statement Types

- Data Definition (Create Table)
- Data Manipulation (Insert, Update, Delete)
- Data Retrieval (Select)
- Multi-Table Selection (Joins)
- Advanced Summary (Group By)
- View Definition (Create View)

Create Table

Create table Example (
ItemNum int,
ItemName text); ← Don't Forget The
SemiColon

Verify With .tables

Drop Table Example; ← This Will Get Rid of It

Data Manipulation: Insert

Insert into Example (ItemNum, ItemName)
Values (1, "First");

Or

Insert into Example Values (2, "Second");
Insert into Example Values (3, "Third");
Insert into Example Values (4, "Fourth");
Insert into Example Values (5, "Fifth");
We Need a Few Rows to Work With

Data Retrieval: Select

Select * from Example;

You Can Make It Easier to Read By Entering:

.separator "\t" "\n"
 ↑ ↑
 | |
 | |
Column Line
Separator Separator

Data Manipulation: Delete

Delete from Example
Where ItemNum = 1;

Verify With Select

Delete from Example
With No Criteria Will Delete All Rows

Data Manipulation: Update

Update Example Set ItemName = "Third"
Where ItemNum = 2;

Verify With Select

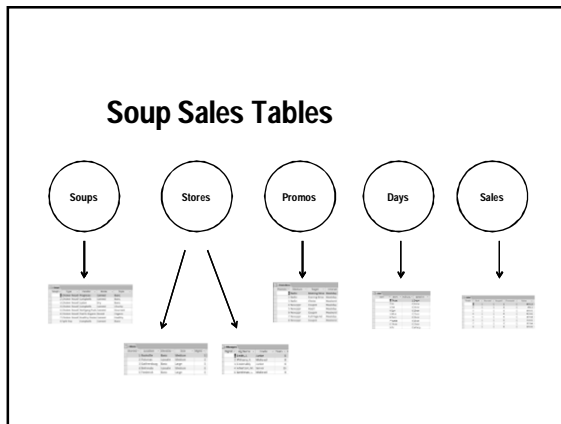
Update Example Set ItemName = "Fourth";
With No Criteria Will Update All Rows

BTW, This Works With Access Too

- Invoke MS Access
- Use Create Query to Get to the SQL Interface
- Create a Table
- Insert a Row
- Update the Row
- Verify the Updates
- And It Will Work With Any Relational Database

Now Some More Realistic Data

- We Have Some Data Soup Sales That We Wish to Analyze. We Will:
 - Write Create Table Statements for Them
 - Define Them to the Database
 - Load Data Into Them
 - Write SQL Queries to Summarize the Data
- This is an Simple Example of What is Called *Data Wrangling*



Now, We Will...

- Define the Tables Using SQL
- Use That SQL to Create the Tables in SQLite3
- Use SQL to Populate and Update the Tables
- See How to Bulk Load Data Into Tables
- And Write Some Simple SQL Queries Using Those Tables to Answer Some Questions

Creating and Populating Days

```
create table Days (
  DoY int not null,
  DoW text,
  Holiday int,
  Weather text);
```

```
.import Days.txt Days
```

```
Verify table with .tables
Verify rows with select
```

Creating and Populating Managers

```
create table Managers (  
  MgrId int not null,  
  MgrName text,  
  Grade text,  
  Years int);  
  
.import Managers.txt Managers  
  
Verify table with .tables  
Verify rows with select
```

Creating and Populating Promos

```
create table Promotions (  
  PromId int not null,  
  Medium text,  
  Target text,  
  Interval text);  
  
.import Promotions.txt Promotions  
  
Verify table with .tables  
Verify rows with select
```

Creating and Populating Soups

```
create table Soups (  
  SoupId int not null,  
  Type text,  
  Vendor text,  
  Mode text,  
  Style text);  
  
.import Soups.txt Soups  
  
Verify table with .tables  
Verify rows with select
```

Creating and Populating Stores

```
create table Stores (  
  StoreId int not null,  
  Location text,  
  Size text,  
  Elevation text,  
  MgrId int);  
  
.import Stores.txt Stores  
  
Verify table with .tables  
Verify rows with select
```

Creating and Populating Sales

```
create table sales (  
  TrxId int not null,  
  DoY int,  
  StoreID int,  
  Soupld int,  
  Promold int,  
  Sales number);  
  
.import Sales.txt Sales  
  
Verify table with .tables  
Verify rows with select
```

Bulk Load

- We Can Put All the Create Table Statements and the .import Statements into a file called tabledefs.txt
- We Can Define the Database and Load the Tables With One Command
- .read tabledefs.txt (type .echo on to echo cmds)
- Verify With .tables

Now That We Have Data....

Kinds of SQL Queries

- Single Table Queries – This Week
- Multi-Table Queries – Next Week
 - Nested Queries
 - Joins (Next Semester)
- Aggregation Queries – Next Week
 - Single Table Aggregation
 - Multi-Table Aggregation (Next Semester)
 - Advanced Queries (Next Semester)

Basic Select – List a Table

Select * From Days;
Type .headers on for Column Headers

Limit and Offset

```
Select * From Days  
Limit 10 Offset 5;
```

Shows Rows 6 Through 15

Select Some of the Columns

```
Select DoY, DoW  
From Days;
```

Select Some of the Rows

```
Select type from Soups  
Where Vendor = "Campbells"
```

Compound Conditions

Select type from Soups
Where Vendor = "Campbells"
Or Vendor = "Progresso";

Select type from Soups
Where Vendor = "Campbells"
And Mode = "Canned";

Truth Tables

Operator	Value 1	Value 2	Result
And	True	True	True
And	True	False	False
And	False	True	False
And	False	False	False
Or	True	True	True
Or	True	False	True
Or	False	True	True
Or	False	False	False

This is the Same as For Boolean Expressions

Select: Removing Duplicates

Select distinct(type) from Soups;

Between Example

Select Sales from Sales
Where sales between 200 and 259;

Like Example

Select type from Soups
Where type like "Chicken%";

Ordering Results

Select Years, MgrName from Managers
Order by Years;

Queries Using Missing Values

Select Soupld from Soups
Where Type is Null;

Truth Tables for 3 Value

3VL And	True	False	Null
True	True	False	Null
False	False	False	False
Null	Null	False	Null

3VL Or	True	False	Null
True	True	True	True
False	True	False	Null
Null	True	Null	Null

I Am Working on This
So Don't Worry About It

Compound Conditions W/Nulls

Select type from Soups
Where Vendor = "Campbells"
Or Vendor = "Progresso";

Select type from Soups
Where Vendor = "Campbells"
And Mode = "Canned";

I Am Working on Some Examples For
This So, Don't Worry About it For Now

In Example

Select Weather from Days
Where DoY in (1,8,9,12);

Not In

Select Distinct(Weather) from Days
Where DoY Not in (1,8,9,12);

Summary

- Introduction to SQLite3 and Demo
- Creating Tables in SQL
- SQL Data Manipulation Statements
- Defining the SoupSales Tables
- Some Simple SQL Queries
